

# The Privacy Externality Game: Strategic Budget Choices in Sampling-Based Individual Differential Privacy

Anonymous Submission

**Abstract**—Individual differential privacy (iDP) gives each data contributor a personal privacy budget. The currently dominant realisation—sampling-based iDP—turns each contributor’s chosen budget into a per-sample inclusion rate, with a single noise multiplier shared across the dataset. We show that this realisation creates a privacy externality: lowering one’s own budget mechanically raises the realised membership-inference (MIA) advantage on every other data point, even though every individual  $(\varepsilon_i, \delta)$ -iDP guarantee continues to hold. We model the resulting strategic situation as a normal-form *Privacy Externality Game* on a continuous action space, with each player’s utility trading off a shared model-quality term  $V(\sigma)$  against their own realised MIA advantage  $A_i(\varepsilon)$ . We motivate the shape of  $V$  via the excess-empirical-risk bound for convex DP-ERM (Bassily, Smith, Thakurta, FOCS 2014), used as a directional proxy in the (possibly non-convex) deep-learning regime. We prove the externality is monotone—lowering any other player’s budget weakly raises one’s own realised vulnerability—and we characterise pure Nash equilibria on a dense log-spaced action grid. In the symmetric 2-player setting the pure NEs are *asymmetric*: at any symmetric profile both players have a deviation incentive in the same direction, and the equilibrium settles with one player carrying a higher realised  $A$  in exchange for a higher  $V$ . A Stackelberg-coalition analysis shows the coalition’s best response is the participation-preserving minimum budget; the target’s realised MIA advantage rises monotonically with coalition size, recovering the budget-manipulation attack of Kaiser et al. (SaTML 2026) as the rational best response of any coordinated set of data holders.

**Index Terms**—Differential privacy, individual differential privacy, game theory, Nash equilibrium, collusion, membership inference.

## I. INTRODUCTION

Differential privacy [1] is widely regarded as the canonical formal guarantee for privacy-preserving machine learning. In its classical form, a single privacy budget  $(\varepsilon, \delta)$  is shared by every record, and every record receives the same protection. Yet privacy preferences are heterogeneous: a clinical patient and a casual user contributing data to the same model often place very different values on their privacy. *Individual differential privacy* (iDP) replaces the single  $(\varepsilon, \delta)$  guarantee with a per-record contract  $(\varepsilon_i, \delta)$ , and several mechanisms have been proposed to deliver such per-record contracts in deep learning [2]–[4]. The most widely-adopted realisation in DP-SGD is sampling-based iDP, due to [2]: heterogeneous budgets are achieved by calibrating each record’s Poisson inclusion probability  $p_i$  against a single shared noise multiplier  $\sigma$  so that, for fixed batch size and iteration count, every group of records satisfies its own  $(\varepsilon_g, \delta)$ -DP bound.

We focus on a fundamental, overlooked feature of this realisation. Because every record shares the same noise multiplier and a common expected batch size, each record’s sampling rate is mechanically coupled to every other record’s chosen budget. Lowering one’s own  $\varepsilon_i$  forces other contributors’ sampling rates upward to keep the expected batch size fixed—which, as recent empirical work has shown [5], inflates their realised membership-inference (MIA) advantage even though every formal  $(\varepsilon_g, \delta)$ -iDP guarantee continues to hold. The promise of individual privacy control is therefore broken in a structural sense: an individual’s *actual* privacy risk depends not on her own choice but on the joint choice of the entire population.

a) *Our viewpoint: this is a game.*: The natural framing for any choice problem whose outcome is determined by the choices of multiple self-interested parties is game theory. [6] introduced a two-player game on the privacy-accuracy plane (the *Price of Privacy*), but they were studying a mechanism in which a player’s privacy choice affected only her *own* accuracy. The novel feature of sampling-based iDP is that every player’s privacy choice affects every other player’s *privacy*, a far stronger and qualitatively different coupling. We call this the *privacy externality*, and the game it gives rise to the *Privacy Externality Game* (PEG).

b) *Contributions.*: This paper makes the following contributions.

- 1) **Formalisation.** We define the  $n$ -player Privacy Externality Game on a continuous action space. Each player chooses a per-group budget  $\varepsilon_i \in [\varepsilon^{\min}, \varepsilon^{\max}]$ ; the sampling-based iDP mechanism then determines the noise multiplier  $\sigma$  and the per-group sampling rates  $p_g$ , and each player’s utility is the canonical scalar  $u_i = B_i V(\sigma) - C_i A_i(\varepsilon)$  trading off model quality against the realised MIA advantage on the player’s own data. We motivate the functional form of  $V$  via the excess-empirical-risk bound of [7] for convex DP-ERM, noting the directional caveat as we use it in the (potentially non-convex) deep-learning regime.
- 2) **Externality lemma.** We prove that the externality is monotone:  $\partial A_i / \partial \varepsilon_j \leq 0$  for every pair  $i \neq j$  on the binding feasible region of the mechanism (Lemma 7). Lowering any other player’s budget never reduces my realised vulnerability.
- 3) **Equilibrium characterisation.** We show numerically on a dense log-spaced action grid that the symmetric 2-

player PEG admits only *asymmetric* pure Nash equilibria: at any symmetric profile, both players have an incentive to deviate in the same direction, and the equilibrium settles with one player carrying a higher realised  $A$  in exchange for a higher  $V$ . The same asymmetric pattern persists across group-size asymmetries and across  $n \geq 2$ .

- 4) **Coalition Stackelberg recovers the attack.** A coalition of data holders maximising the target's realised MIA advantage picks the participation-preserving minimum  $\varepsilon^*$ , and the target's  $A_t$  is monotone-increasing in coalition size  $k$  (Proposition 9). This exactly recovers the budget-manipulation attack of [5] as the rational best response of any coordinated set of data holders. In our  $n = 5$  canonical setting the target's  $A_t$  exceeds 0.7 at  $k = 4$ .
- 5) **Empirical evaluation.** We numerically evaluate the game on a range of small but representative scenarios mirroring [5]. We report equilibrium budgets, realised MIA advantages, and the coalition-attack curve.

c) *Implication.*: The prior attack literature on iDP frames the excess vulnerability as a *vulnerability*—something a malicious party exploits. We argue the opposite. Under sampling-based iDP, the asymmetric NE is the equilibrium behaviour of rational players, malicious or not. Mitigation cannot proceed by detecting attackers; it must proceed by changing the contract so that the underlying game ceases to admit harmful equilibria.

The paper is organised as follows. Section II reviews differential privacy, sampling-based iDP, MIA advantage, and basic game theory. Section III formalises the Privacy Externality Game. Section IV proves the externality lemma and the coalition-attack recovery and characterises equilibrium structure. Section V reports the numerical evaluation. Section VI reviews related work and Section VII concludes.

## II. BACKGROUND

### A. Differential privacy and Gaussian mechanisms

We use the standard definition of approximate differential privacy [1]: a randomised mechanism  $\mathcal{M}: \mathcal{Z}^* \rightarrow \mathcal{O}$  satisfies  $(\varepsilon, \delta)$ -differential privacy  $((\varepsilon, \delta)$ -DP) if, for any pair of adjacent databases  $D \sim D'$  (differing on a single record) and every measurable  $S \subseteq \mathcal{O}$ ,

$$\Pr[\mathcal{M}(D) \in S] \leq e^\varepsilon \Pr[\mathcal{M}(D') \in S] + \delta.$$

A single  $(\varepsilon, \delta)$ -pair is only a single point of the underlying *privacy profile*  $\{(\varepsilon, \delta(\varepsilon)) : \varepsilon \geq 0\}$  [8], [9], and mechanisms calibrated to the same point can have substantially different profiles and so substantially different realised adversarial risk. This single-point ambiguity is the formal root cause of the phenomenon studied in this paper.

The *Gaussian mechanism*,  $\mathcal{M}(D) = f(D) + \mathcal{N}(0, \sigma^2 I)$  for a query  $f$  of  $L_2$  sensitivity  $\Delta f$ , satisfies  $(\varepsilon, \delta)$ -DP for all  $(\varepsilon, \delta)$  such that  $\Phi(\Delta f/2\sigma - \varepsilon\sigma/\Delta f) - e^\varepsilon \Phi(-\Delta f/2\sigma - \varepsilon\sigma/\Delta f) \leq \delta$ , where  $\Phi$  is the standard-normal CDF. For training neural networks under DP, the standard procedure is DP-SGD with Poisson subsampling [10], in which each data point is included in each minibatch independently with probability  $p$ , providing the well-known *subsampling amplification* of DP.

### B. Individual differential privacy and the sampling-based realisation

*Individual differential privacy* [3] relaxes the uniform  $(\varepsilon, \delta)$  guarantee to a per-record contract  $(\varepsilon_i, \delta)$ : the resulting algorithm must satisfy  $(\varepsilon_i, \delta)$ -DP with respect to record  $i$  for every  $i$ . Boenisch et al. [2] translate per-record contracts into per-record Poisson sampling probabilities while keeping a single noise multiplier  $\sigma$  and a fixed expected batch size, which is the configuration we study.

**Definition 1** (Sampling-based iDP-SGD [2]). Let the dataset of size  $M$  partition into  $G$  *privacy groups*  $\{G_1, \dots, G_G\}$  with sizes  $|G_g|$ . Each group has a target  $(\varepsilon_g, \delta)$ -DP budget. The mechanism trains by Poisson-subsampled DP-SGD for  $I$  iterations, where the noise multiplier  $\sigma$  and the per-group sampling rates  $p_1, \dots, p_G$  are chosen such that

(C1) for every  $g$ , the  $I$ -fold  $(p_g, \sigma)$ -subsampled

Gaussian satisfies  $(\varepsilon_g, \delta)$ -DP at equality.

$$(C2) \quad \sum_{g=1}^G \frac{|G_g|}{M} p_g = \frac{E_b}{M}, \quad (1)$$

where  $E_b$  is the desired expected batch size. We refer to  $(\sigma, p_1, \dots, p_G)$  as the *mechanism solution* for the budget vector  $\varepsilon = (\varepsilon_1, \dots, \varepsilon_G)$ .

The single shared  $\sigma$  in (C1) is the source of the privacy externality. For any candidate  $\sigma$ , the largest sampling rate that keeps group  $g$  at its budget is a monotone-increasing function of  $\sigma$  (more noise  $\Rightarrow$  more room for sampling); the batch-size constraint (C2) then determines  $\sigma$  as the unique value at which the per-group rates sum to  $E_b$ . As a corollary, the rate  $p_g$  that group  $g$  obtains is determined by the entire vector  $\varepsilon$ , not by  $\varepsilon_g$  alone.

### C. MIA advantage and the privacy profile

For a privacy-preserving mechanism  $\mathcal{M}$  with trade-off function  $T$  (equivalently, an  $f$ -DP function [11]), the membership inference advantage of the optimal Bayes-aware adversary is [5]

$$\text{Adv}(\mathcal{M}) = \sup_{\alpha \in [0,1]} (1 - \alpha - T(\alpha)), \quad (2)$$

i.e. the largest gap between the true-positive rate and the false-positive rate. For the  $I$ -fold  $(p, \sigma)$ -subsampled Gaussian,  $T$  can be computed numerically by privacy-loss-distribution (PLD) accounting [12], which is what we use throughout. Crucially,  $\text{Adv}$  is determined by the *full privacy profile*, not by any single  $(\varepsilon, \delta)$  point.

### D. Game-theoretic terminology

A finite normal-form game is a tuple  $\langle N, \Sigma, U \rangle$  with  $N = \{1, \dots, n\}$  players, action sets  $\Sigma = \{S_1, \dots, S_n\}$ , and payoff functions  $U = \{u_1, \dots, u_n\}$ . A *Nash equilibrium* (NE) is an action profile  $\mathbf{s}^*$  such that no player profits from a unilateral deviation:  $u_i(\mathbf{s}_i^*, \mathbf{s}_{-i}^*) \geq u_i(\mathbf{s}_i, \mathbf{s}_{-i}^*)$  for every  $\mathbf{s}_i \in S_i$  [13]. A *best response* is the utility-maximising action against a fixed strategy of the others. The *price of anarchy* [14] compares the

welfare of the worst NE to the social optimum. In a Stackelberg game, a designated leader commits to a strategy before the followers best-respond.

### III. THE PRIVACY EXTERNALITY GAME

We formalise the strategic situation faced by data contributors in sampling-based iDP as a normal-form game on a continuous action space. The formalisation is a direct extension of [6]’s two-player Collaborative-Learning game, with the central modification that the *privacy* cost is no longer separable across players.

#### A. Players, actions and mechanism

**Definition 2** (Privacy Externality Game). Fix a common  $\delta > 0$ , an expected batch size  $E_b$ , an iteration count  $I$ , and a partition of the dataset into  $n$  groups of sizes  $|G_1|, \dots, |G_n|$ . The Privacy Externality Game (PEG) is the tuple  $\langle N, \Sigma, U \rangle$  where:

- $N = \{1, \dots, n\}$  is the set of players. We adopt the FL-silo granularity of [5]: player  $i$  is the entity controlling group  $i$  (a silo). For a single-silo single-budget formulation each player picks a scalar  $\varepsilon_i$ ; the framework extends to per-record distributions inside a silo, which we discuss in Section VII.
- Each player’s action set is  $\Sigma_i = [\varepsilon^{\min}, \varepsilon^{\max}] \subset \mathbb{R}_{>0}$ .
- Given the action profile  $\varepsilon \in \prod_i \Sigma_i$ , the mechanism solution  $(\sigma(\varepsilon), p_1(\varepsilon), \dots, p_n(\varepsilon))$  is determined by Definition 1.
- Each player’s utility is

$$u_i(\varepsilon) = B_i \cdot V(\sigma(\varepsilon)) - C_i \cdot A_i(\varepsilon), \quad (3)$$

with  $B_i, C_i > 0$ .

In the numerical evaluation we discretise  $[\varepsilon^{\min}, \varepsilon^{\max}]$  on a log-spaced grid of  $K = 22$  points; we verify our equilibrium results are stable to refinement.

#### B. The two ingredients of utility

a) *Model-utility value*  $V(\sigma)$ .: We model model utility as a monotone-decreasing scalar function of the mechanism’s noise multiplier  $\sigma$ . We use

$$V(\sigma) = \frac{1}{1 + \sigma^2}. \quad (4)$$

*Frame.* The convex-DP-ERM excess-empirical-risk bound of [7] (FOCS 2014) is, for a Lipschitz strongly-convex loss,  $O(d \log(1/\delta)/(n\varepsilon)^2)$ , which translates under the moments accountant of [10] (CCS 2016) into a  $\sigma^2/n^2$  scaling of the excess risk. *Validity in our setting.* The bound is proven in the convex setting; deep-learning DP-SGD is generally non-convex and the bound does not apply directly. We use  $V(\sigma) = 1/(1 + \sigma^2)$  as a *directional* proxy. The empirical literature on DP-SGD in deep learning [2], [10], [15] reports monotonic degradation of test accuracy with  $\sigma$ , supporting the qualitative form. None of our game-theoretic results depend on the precise shape of  $V$ : only that it is continuous, monotone-decreasing in  $\sigma$ , and bounded.

b) *Realised MIA advantage*  $A_i(\varepsilon)$ .: We use the tight analytical MIA advantage of the  $I$ -fold  $(p_i, \sigma)$ -subsampled Gaussian as defined in (2). This is the same advantage measure used by [5]; our choice transfers their empirical measurement to a strategic setting unchanged. The functional dependence  $A_i(\varepsilon)$  runs through  $(p_i, \sigma)$ , both of which are determined by the entire vector  $\varepsilon$  via the mechanism’s (C1) and (C2).

#### C. Decomposition into contracted and excess cost

The privacy cost  $A_i(\varepsilon)$  in (3) can be decomposed into a contracted (separable) part and an excess (coupled) part that makes the externality explicit. Let  $\bar{c}(\varepsilon_i)$  denote the MIA advantage of a reference, non-interdependent  $(\varepsilon_i, \delta)$ -DP mechanism (e.g., the post-processing-tight  $(\varepsilon_i, \delta)$ -DP advantage upper bound  $(e^{\varepsilon_i} - 1 + 2\delta)/(e^{\varepsilon_i} + 1)$  of [9], [16]). The realised cost is

$$A_i(\varepsilon) = \bar{c}(\varepsilon_i) - \text{slack}_i(\varepsilon), \quad (5)$$

where the *slack*  $\text{slack}_i(\varepsilon) := \bar{c}(\varepsilon_i) - A_i(\varepsilon) \geq 0$  is the amount by which the realised advantage falls below the worst-case-bound at  $(\varepsilon_i, \delta)$ . In sampling-based iDP this slack is a function of the entire  $\varepsilon$  through the mechanism’s coupling. *The externality reduces the slack*: when others lower their  $\varepsilon_j$ ,  $\text{slack}_i$  shrinks (player  $i$  moves closer to its  $(\varepsilon_i, \delta)$ -DP worst case). This is the formal quantity that the proposed  $(\varepsilon_i, \delta, \Delta)$ -iDP contract [5], [9] would cap: it would require  $\text{slack}_i$  to differ from a baseline reference by at most  $\Delta$ . In this sense, *the  $\Delta$ -bound is exactly the device that neutralises the externality*.

We will work directly with  $A_i$  in the utility (3) for analytical simplicity; the decomposition matters as a conceptual bridge to PoP [6] (whose cost  $c(p_i)$  was separable, hence  $\text{slack}_i$  did not exist) and to the iDP literature [5] (where  $\text{slack}_i$  is the new externality term).

#### D. Vulnerability function and feasibility

**Definition 3** (Vulnerability function). The *vulnerability function* of player  $i$  is  $A_i: \prod_i \Sigma_i \rightarrow [0, 1]$ ,  $\varepsilon \mapsto \text{Adv}(\mathcal{M}_{p_i(\varepsilon), \sigma(\varepsilon)}, I)$ .

**Assumption 4** (Participation). A profile  $\varepsilon$  is *feasible* if and only if the mechanism solution satisfies  $p_i(\varepsilon) > 0$  for every player  $i$ . Profiles with  $p_i = 0$  correspond to a player whose data never enters training; their iDP contract is satisfied trivially. We exclude such profiles from the game; equivalently, the action space  $\Sigma_i$  is implicitly restricted at each move to the participating region.

#### E. Coalition variant

**Definition 5** ( $k$ -coalition Stackelberg PEG). Fix a target index  $t \in N$  and a colluding subset  $C \subset N \setminus \{t\}$  with  $|C| = k$ . The  $k$ -coalition Stackelberg PEG is the two-stage game in which:

- 1) Stage 1 (leader): the coalition jointly commits to an action vector  $\varepsilon_C$  to maximise the augmented coalition utility

$$\sum_{i \in C} u_i(\varepsilon) + \lambda A_t(\varepsilon), \quad (6)$$

where  $\lambda \geq 0$  is the *malice weight* on the target's vulnerability [6].  $\lambda = 0$  recovers the benign PEG;  $\lambda \rightarrow \infty$  recovers the budget-manipulation attack of [5].

- 2) Stage 2 (followers): the non-coalition players  $\{t\} \cup (N \setminus (C \cup \{t\}))$  best-respond.

#### F. Equilibrium concepts

We use the standard Nash equilibrium concept: an  $\varepsilon^*$  in the feasible region such that  $u_i(\varepsilon_i^*, \varepsilon_{-i}^*) \geq u_i(\varepsilon_i, \varepsilon_{-i}^*)$  for every player  $i$  and every feasible deviation  $\varepsilon_i$ . Existence of pure NE on the continuous action space is not automatic; we report pure NE found by brute search on a fine log-spaced grid and verify the result is stable under grid refinement. Where no pure NE exists on the dense grid we report the empirical mixed strategy reached by fictitious play [17].

*Remark 6* (Awareness as information, not utility). A useful distinction is between players who can *compute*  $A_i(\varepsilon)$  given the public mechanism's parameters and players who can only access their nominal  $(\varepsilon_i, \delta)$ -iDP contract. Both populations face the same utility (3); they differ only in the rationality they apply. A *nominal-information* ("naive") player who treats the contracted bound  $\bar{c}(\varepsilon_i)$  as a separable cost plays the PoP game of [6] unchanged. An *aware* player who computes  $A_i(\varepsilon)$  from the mechanism's parameters plays the PEG. We will return to this distinction in Section VII.

### IV. EQUILIBRIUM ANALYSIS

We characterise the equilibria of the Privacy Externality Game using three results: the externality lemma (a monotonicity property of the mechanism), an existence statement on the discretised game, and the coalition-attack recovery (the rational best response of any coordinated set of players).

#### A. Monotonicity of the externality

**Lemma 7** (Privacy externality on the feasible region). *Consider the sampling-based iDP mechanism of Theorem 1 on a feasible profile  $\varepsilon$  at which constraints (C1) bind for every group and (C2) binds at  $E_b$ . Let  $\sigma$  and  $p_g$  denote the mechanism's solution. Then for every pair  $i \neq j$ ,*

$$\frac{\partial \sigma}{\partial \varepsilon_j} < 0 \quad \text{and} \quad \frac{\partial p_i}{\partial \varepsilon_j} < 0. \quad (7)$$

*That is, when player  $j$  loosens its budget, the mechanism uses less noise ( $\sigma$  falls) and reduces group  $i$ 's sampling rate ( $p_i$  falls). Empirically (Fig. 4), the realised MIA advantage on group  $i$  also falls:  $\partial A_i / \partial \varepsilon_j \leq 0$ .*

*Proof.* Use the Rényi-DP form of (C1) of [18]: at order  $\alpha$ , the RDP of an  $I$ -fold  $(p_g, \sigma)$ -subsampled Gaussian is at most  $\rho_g \approx 2I\alpha p_g^2 / \sigma^2$  (tight up to higher-order terms for  $p_g \ll 1$ ). Setting this equal to the target RDP  $\rho_g(\varepsilon_g, \delta)$  (which is monotone-increasing in  $\varepsilon_g$  because  $\rho = \varepsilon - \log(1/\delta)/(\alpha - 1)$ ) for the binding RDP-to-DP conversion at any chosen order  $\alpha > 1$ ) gives the per-group sampling rate

$$p_g(\sigma, \varepsilon_g) = \frac{\sigma}{\sqrt{2I\alpha}} \sqrt{\rho_g(\varepsilon_g, \delta)}.$$

Substituting into (C2) yields

$$\sigma \cdot \sum_g |G_g| \sqrt{\rho_g(\varepsilon_g, \delta)} = E_b \sqrt{2I\alpha}. \quad (8)$$

Implicit differentiation in  $\varepsilon_j$  gives  $\partial \sigma / \partial \varepsilon_j = -\sigma |G_j| (\partial \sqrt{\rho_j} / \partial \varepsilon_j) / [\sum_g |G_g| \sqrt{\rho_g}] < 0$ , since  $\rho_j$  is increasing in  $\varepsilon_j$ . From the closed-form expression for  $p_i$  at fixed  $\varepsilon_i$ ,  $\partial p_i / \partial \varepsilon_j = (\partial \sigma / \partial \varepsilon_j) \cdot \sqrt{\rho_i(\varepsilon_i)} / \sqrt{2I\alpha} < 0$ .

The induced sign of  $\partial A_i / \partial \varepsilon_j$  is not determined by  $\partial p_i / \partial \varepsilon_j$  and  $\partial \sigma / \partial \varepsilon_j$  alone: the MIA advantage of the  $I$ -fold  $(p_i, \sigma)$ -subsampled Gaussian depends on the joint shape of the privacy profile, and the chain  $\partial A_i / \partial p_i$  and  $\partial A_i / \partial \sigma$  have opposite signs (more sampling exposes more; more noise hides more). However, we verify numerically in Fig. 4 (and consistent with Figure 3 of [5]) that  $\partial A_i / \partial \varepsilon_j \leq 0$  throughout the feasible region of our scenarios: the  $p_i$ -decrease effect dominates the  $\sigma$ -decrease effect.  $\square$

Theorem 7 formalises the empirical observation of [5]: when any player tightens their iDP budget, the mechanism's noise multiplier  $\sigma$  shifts to keep (C2) binding, and every other player's realised vulnerability  $A_i$  rises. The externality is *monotone* in every other player's budget.

In the decomposition of Section III-C, the lemma equivalently states that *the slack  $\bar{c}(\varepsilon_i) - A_i(\varepsilon)$  shrinks in every other player's  $\varepsilon_j$* . In words, the  $(\varepsilon_i, \delta)$ -DP worst-case ceiling  $\bar{c}(\varepsilon_i)$  is unchanged (it depends only on  $\varepsilon_i$ ), but the realised advantage  $A_i$  climbs toward that ceiling as others tighten their budgets. This is the quantity that the proposed  $(\varepsilon_i, \delta, \Delta)$ -iDP contract [5], [9] would cap, and our Lemma motivates such a cap formally: without it,  $\text{slack}_i$  has no floor and the externality can drive any player's realised advantage arbitrarily close to the worst case.

*a) Minority and majority intuition.:* The externality has an asymmetric *group-size* structure. Because (C2) enforces a fixed expected batch size, a player with small  $|G_i|$  but small  $\varepsilon_i$  contributes little to the batch and forces the larger- $|G_j|$  players' sampling rates upward. The minority therefore *exports* its privacy demand onto the majority. This is the formal content of the empirical observation in [5] that "the proportion of low-budget players matters as much as their  $\varepsilon$ ".

#### B. Existence and structure of pure Nash equilibria

**Proposition 8** (NE existence on the discretised game). *The discretised PEG on the action grid  $\mathcal{A}_K = \{\varepsilon^{\min}, \dots, \varepsilon^{\max}\} \subset [\varepsilon^{\min}, \varepsilon^{\max}]$  with utilities (3) is a finite normal-form game restricted to the (closed) feasible set. By the Nash existence theorem [13], every such game admits at least one mixed-strategy NE. Pure NE may or may not exist, depending on  $(B_i, C_i, |G_i|, V, K)$ .*

We do not claim existence on the continuous action space without further quasi-concavity assumptions; in particular, our utility (3) is not generally concave in  $\varepsilon_i$  because  $A_i$  is the trade-off function of a subsampled Gaussian, whose dependence on  $\varepsilon_i$  through  $(\sigma, p_i)$  is non-trivial. We instead report pure NE

found by exhaustive search on a dense log-spaced grid and verify that they survive grid refinement.

a) *What the externality lemma implies for NE structure.*

Lemma 7 says lowering another player’s budget weakly raises one’s own realised vulnerability. If players’ utilities place non-trivial weight on  $A_i$  ( $C_i > 0$ ), each player has an incentive to *shift the externality onto others* by lowering their own budget—an incentive that is intrinsic to the mechanism. The opposing incentive comes from  $V(\sigma)$ : lowering own  $\varepsilon$  pushes  $\sigma$  upward, lowering the model utility for everyone (including oneself). The equilibrium balances these two incentives.

We observe in our experiments (Section V) that the resulting balance is typically *asymmetric*: in the symmetric 2-player setting, pure NEs sit off the diagonal of the action grid. The intuition is the externality lemma’s converse: at any symmetric profile  $(\varepsilon, \varepsilon)$ , both players face an identical incentive to deviate; since both deviations cannot be unilaterally absorbed, the equilibrium is at an asymmetric profile where one player carries a higher  $A$  in exchange for a higher  $V$ .

### C. Coalition Stackelberg and the attack recovery

**Proposition 9** (Coalition best response recovers the attack). *Fix a target  $t$  and a coalition  $C$  of size  $k$ , with  $\lambda \rightarrow \infty$  in Definition 5 (the coalition cares only about the target’s vulnerability). The coalition’s best response is  $\varepsilon_i = \varepsilon_k^*$  for every  $i \in C$ , where  $\varepsilon_k^*$  is the smallest admissible budget that keeps the coalition’s mechanism solution feasible (i.e.,  $p_i > 0$  for every  $i \in C$ ). Furthermore, the target’s MIA advantage  $A_t$  is non-decreasing in  $k$ .*

*Proof sketch.* By Theorem 7,  $\partial A_t / \partial \varepsilon_i \leq 0$  for every  $i \in C$  on the feasible region. Setting every  $\varepsilon_i$  as low as the feasibility region admits therefore maximises  $A_t$ . The participation-preservation qualifier is necessary because at extremely low  $\varepsilon_i$  the mechanism may set  $p_i = 0$  (the colluder effectively drops out, removing its constraint from (C2) and so its externality on the target); the coalition’s optimum is thus the infimum of feasible  $\varepsilon_i$ . Monotonicity in  $k$  is the same lemma applied to adding a member to  $C$ : doing so strictly enlarges the set of budgets that pull  $\sigma$  downward, raising  $A_t$ .  $\square$

Proposition 9 establishes that the budget-manipulation attack of [5] is precisely the *rational* best response of any coordinated subset of data holders, regardless of malicious intent. Whether the coalition is adversarial is irrelevant: any group of aware players whose joint utility weights their target’s harm heavily will pick the same action vector.

### D. An aside on awareness

Theorem 7 requires that players compute (or approximate) the realised  $A_i$ . A natural question is: what about a player whose only information is the formal  $(\varepsilon_i, \delta)$ -iDP contract? Such a player treats  $\varepsilon_i$  as a direct privacy proxy and plays the budget that maximises  $V(\sigma) - g(\varepsilon_i)$  for some monotone  $g$ . Their best response does not depend on  $\varepsilon_{-i}$  at all (the externality is invisible), so their action is mechanically the smallest  $\varepsilon_i$  that does not break their own  $V$ . This is the *rational* action of

TABLE I: Overview of the four numerical experiments.

|    | Question                              | Setup                                                                 | Finding                                                                                 |
|----|---------------------------------------|-----------------------------------------------------------------------|-----------------------------------------------------------------------------------------|
| C1 | Symmetric equilibrium?                | 2 equal-size groups, full $\varepsilon$ -grid sweep                   | 8 NEs, all asymmetric (Fig. 1)                                                          |
| C2 | Effect of size asymmetry              | 2 groups, size ratio $r \in [0.1, 10]$                                | Minority plays low $\varepsilon$ , majority absorbs $A$ (Fig. 2)                        |
| C4 | How bad under coordination?           | $n=5$ , target $\varepsilon_t=32$ , coalition of size $k$             | $A_t$ monotone in $k$ ; +65% at $k=4$ (Fig. 3)                                          |
| C5 | Where does the externality come from? | 2-group, vary $\varepsilon_O$ and proportion at fixed $\varepsilon_T$ | $A_t$ monotone in $\varepsilon_O$ , crossover at $\varepsilon_O=\varepsilon_T$ (Fig. 4) |

a naive player, and it is precisely the configuration in which the attack of [5] succeeds against them: by Lemma 7, naive players’ choices feed the externality on everyone else without their realising. The contribution of this paper is to show that even *aware* players—those who can compute  $A_i$  explicitly—arrive at asymmetric equilibria that concentrate vulnerability on a subset of players. The externality cannot be undone by individual rationality alone.

## V. NUMERICAL EVALUATION

We evaluate the PEG numerically across four experiments. Each isolates one distinct question that the framework raises. Table I summarises what each one asks, what it controls for, and what it answers.

### A. Common setup

All experiments are analytical—no model training—and follow the same mechanism-solver and advantage-computation pipeline.

a) *Mechanism solver.* For each action profile  $\varepsilon$  we compute the mechanism solution  $(\sigma, p_1, \dots, p_n)$  by bisecting on  $\sigma$  until the batch-size constraint (C2) is met, with each per-group sampling rate chosen to make the (C1) RDP constraint bind [18]. Profiles in which any  $p_i$  falls below a small participation threshold are flagged as infeasible (Theorem 4). We then compute each realised  $A_i$  exactly from the privacy profile of the  $I$ -fold  $(p_i, \sigma)$ -subsampled Gaussian using PLD accounting [12], recovering the formula of (2).

b) *Action space.* Unless stated, the action space is the log-spaced grid  $\mathcal{A}_K \subset [0.5, 64]$  with  $K = 22$  points,  $\delta = 10^{-12}$ . We verified on a refined grid ( $K = 44$ ) that the NEs reported are stable to grid refinement, so we report them as representative of the continuous-action equilibrium.

c) *Scenarios.* Unless stated, all experiments use  $n$  symmetric groups of  $|G_g| = 5000$  samples each, expected batch size  $E_b = 128$ , 5 epochs of training,  $B_i = C_i = 1$ , and  $V(\sigma) = 1/(1 + \sigma^2)$ . These choices broadly mirror the empirical iDP scenarios of [2], [5].

TABLE II: Pure Nash equilibria of the symmetric 2-player PEG on the log-spaced action grid  $\mathcal{A}_K = 22$  points in  $[0.5, 64]$ ,  $|G_1| = |G_2| = 5000$ ,  $E_b = 128$ , 5 epochs,  $\delta = 10^{-12}$ ,  $B_i = C_i = 1$ ,  $V(\sigma) = 1/(1 + \sigma^2)$ . Social optimum is the feasible action profile maximising  $u_1 + u_2$ .

|                 | $(\varepsilon_1^*, \varepsilon_2^*)$ | $A_1$ | $A_2$ | $\sigma$ |
|-----------------|--------------------------------------|-------|-------|----------|
| NE <sub>1</sub> | (0.79, 5.04)                         | 0.001 | 0.202 | 1.160    |
| NE <sub>2</sub> | (1.26, 6.35)                         | 0.002 | 0.232 | 1.028    |
| NE <sub>3</sub> | (1.59, 8.00)                         | 0.002 | 0.266 | 0.929    |
| NE <sub>4</sub> | (2.52, 3.17)                         | 0.081 | 0.120 | 1.186    |
| NE <sub>5</sub> | (3.17, 2.52)                         | 0.120 | 0.081 | 1.186    |
| NE <sub>6</sub> | (5.04, 0.79)                         | 0.202 | 0.001 | 1.160    |
| NE <sub>7</sub> | (6.35, 1.26)                         | 0.232 | 0.002 | 1.028    |
| NE <sub>8</sub> | (8.00, 1.59)                         | 0.266 | 0.002 | 0.929    |
| social opt      | (5.04, 25.40)                        | 0.007 | 0.535 | 0.590    |

**B. C1 — Does the symmetric 2-player game admit a symmetric NE?**

*a) Setup.:* Two players, equal group sizes  $|G_1| = |G_2| = 5000$ . Sweep both players'  $\varepsilon$  over the dense log-spaced grid; brute-force search for all pure NEs.

*b) Finding.:* No. The BR correspondences cross at *eight* points, all off the diagonal (Fig. 1, centre; Table II). The pairs (0.79, 5.04), (1.26, 6.35), (1.59, 8.00), (2.52, 3.17) and their swaps are pure NEs; the realised  $A$  at these equilibria ranges from  $(A_1, A_2) = (0.001, 0.202)$  at the most asymmetric pair down to (0.081, 0.120) at the closest-to-symmetric pair. No NE sits on the diagonal.

*c) Implication.:* Even when the game is fully symmetric in players, group sizes, and weights, *the equilibrium breaks symmetry*. The intuition is in Theorem 7: at any symmetric profile both players face an identical incentive to deviate downward in  $\varepsilon$ ; only an off-diagonal profile resolves the deadlock by having one player carry  $V$  and the other carry low  $A$ .

**C. C2 — Who absorbs the externality when groups differ in size?**

*a) Setup.:* Two players, dataset size ratio  $|G_1|/|G_2| \in \{0.1, 0.25, 0.5, 1, 2, 4, 10\}$  at fixed total  $M = 10000$ , fixed batch  $E_b=128$ . Full grid sweep at each ratio; for each ratio we report the pure NE with the largest asymmetry in realised  $A$ .

*b) Finding.:* *The minority always plays the lower  $\varepsilon$  and absorbs less of the realised vulnerability*. Fig. 2: for  $r < 1$  (group 1 smaller) the equilibrium has  $\varepsilon_1^* \in [1.5, 3]$  and  $\varepsilon_2^* \approx 10$ , giving  $A_1^* < 0.01$  and  $A_2^* \in [0.24, 0.32]$ . For  $r > 1$  the pattern flips by symmetry. At  $r=1$  the C1 NE structure is recovered.

*c) Implication.:* The externality has a *group-size sign*. Because (C2) enforces a fixed expected batch, a small-group player who demands strong privacy contributes little to the batch and forces the large-group player's  $p$  upward; the latter therefore carries the realised  $A$ . This is a privacy analogue of a free-riding result: small groups gain disproportionately from the externality and large groups pay disproportionately for it.

**D. C4 — How severe does the attack get under coordination?**

*a) Setup.:* Five symmetric players, all with  $|G_g| = 5000$ . Player 0 is the target with  $\varepsilon_t = 32$ . The remaining  $n - 1 = 4$  players are partitioned into a coalition  $C$  of size  $k$  (varied) and a non-coalition fixed at  $\varepsilon = 32$ . The coalition picks the action that maximises  $A_t$  (Stackelberg leader,  $\lambda \rightarrow \infty$ ).

*b) Finding.:*  $A_t$  is *monotone-increasing in  $k$ , with a non-trivial optimal coalition budget*. Fig. 3: baseline  $A_t = 0.442$  (no coalition) rises to 0.483 at  $k=1$ , 0.531 at  $k=2$ , 0.607 at  $k=3$ , and 0.728 at  $k=4$ . The optimum coalition budget is  $\varepsilon^* \approx 6.8$  for  $k \in \{1, 2, 3\}$  and drops to  $\varepsilon^* \approx 3.2$  at  $k = 4$ : the participation-preserving minimum decreases as the coalition grows because more colluders can absorb the same externality between them while keeping each  $p_i > 0$ .

*c) Implication.:* The budget-manipulation attack of [5] is the rational best response of a coordinated subset of data holders in the PEG. The empirical 62% attack-success rate reported in their work corresponds to the same equilibrium our framework predicts; Fig. 3 traces its severity as a function of the coordinated coalition's size.

**E. C5 — What is the functional form of the externality?**

*a) Setup.:* Two groups, target with budget  $\varepsilon_T \in \{4, 16\}$  fixed, others' budget  $\varepsilon_O$  swept over the log-spaced grid. The other group's proportion of the total dataset is varied over  $\{0.2, 0.4, 0.6, 0.8\}$  at fixed total size  $M = 10000$ . This isolates the target's vulnerability as a pure function of (a) what the others choose and (b) how many of them there are.

*b) Finding.:*  $A_t$  is *monotone-decreasing in  $\varepsilon_O$ ; the proportion-curves cross exactly at  $\varepsilon_O = \varepsilon_T$* . Fig. 4: at  $\varepsilon_T = 4$ ,  $A_t$  ranges from  $\approx 0.20$  (others at  $\varepsilon_O = 0.5$ ) to  $\approx 0$  (others at  $\varepsilon_O = 64$ ). At  $\varepsilon_T = 16$ , the range is 0.55 down to  $\approx 0.05$ . For  $\varepsilon_O < \varepsilon_T$  a larger proportion of others raises  $A_t$ ; for  $\varepsilon_O > \varepsilon_T$  a larger proportion of others lowers  $A_t$ . The crossover at  $\varepsilon_O = \varepsilon_T$  is the regime in which the others' choice and the target's choice contribute identically to  $\sigma$ .

*c) Implication.:* This isolates the externality as a pure mechanism effect, independent of any strategic choice. The shape mirrors Figure 3 of [5], who showed the same dependence empirically through the iDP-SGD trade-off function. C5 confirms that the externality lemma (Theorem 7) is not just empirically monotone but *quantitatively* so: each unit decrease in others'  $\varepsilon$  corresponds to a smooth, predictable increase in the target's realised  $A$ .

## VI. RELATED WORK

*a) Individual differential privacy.:* Classical  $(\varepsilon, \delta)$ -DP [1] applies uniformly to every record. Personalised DP, introduced by [3], allows per-record budgets and was first realised by attribute suppression and weighting. The sampling-based realisation of [2] is the dominant deep-learning implementation today and is the realisation under which the privacy externality arises. [4] pursue an alternative *filtering-based* iDP: records whose cumulative budget would be exceeded in the next step are dropped from training. Filtering does not exhibit the externality (each record's filter is independent), but in deep learning it is




Fig. 1: **C1.** 2-player payoff surface  $u_1(\epsilon_1, \epsilon_2)$  (left), best-response correspondences with pure NEs as black stars (centre), and realised MIA advantage  $A_1$  (right). All eight pure NEs lie off the diagonal.




Fig. 2: **C2.** Equilibrium budget (left) and realised MIA advantage (right) as a function of dataset size ratio  $|G_1|/|G_2|$ . The minority group's  $\epsilon^*$  is always lower than the majority's, and its realised  $A^*$  correspondingly smaller. The flip at  $r=1$  is the boundary at which the minority/majority labels exchange.

impractical due to the catastrophic-forgetting risk reported by [2]. Our PEG is specific to the sampling-based realisation; the externality vanishes in the sensitivity-based realisation, which preserves the noise-to-sensitivity ratio across the dataset.

*b) Excess privacy risk and mechanism comparison.:* Mechanisms calibrated to the same  $(\epsilon, \delta)$  can have very different privacy profiles and thus different realised risks [8], [9], [11]. The  $\Delta$ -divergence of [9] provides a notion of approximate Blackwell-dominance that quantifies the worst-case excess MIA advantage between mechanisms. [5] use this framework to motivate an  $(\epsilon_i, \delta, \Delta)$ -iDP contract that bounds excess vulnerability and demonstrate empirically that without such a bound, adversarial budget choices inflate the MIA susceptibility of 62% of targeted records. Our PEG complements that work: rather than restricting mechanism design we explain why the realised excess vulnerability arises as the equilibrium of self-

interested play under the current iDP contract.

*c) Game-theoretic privacy in collaborative learning.:* [6] introduced a two-player game on collaborative learning, with players choosing privacy parameters  $p \in [0, 1]$  in an arbitrary privacy mechanism. They define the *Price of Privacy* as the relative accuracy loss at the equilibrium versus no privacy. Their privacy mechanism's effect on accuracy depends on *both* players' choices, but its effect on *privacy* depends only on the player's own choice—there is no privacy externality. Our PEG transplants their game-theoretic methodology into a setting where the externality is in the privacy direction, not the accuracy direction. We retain the qualitative contributions of their work (two-player setting, choice of  $\epsilon$ , notion of equilibrium analysis) and extend them to the externality regime.

Earlier work studied related game-theoretic privacy questions in different settings. [19] model linear regression with private






Fig. 3: **C4.** Target’s realised MIA advantage as a function of coalition size  $k$ , with the coalition playing its target-maximising joint action. Annotations show the optimal coalition budget  $\epsilon^*$  at each  $k$ .

features as a non-cooperative game. [20] study non-monetary incentives in data analytics. [21] model the machine-learner-vs-data-perturber conflict as a Stackelberg game. None of these treat the privacy externality of a shared sampling mechanism, which is specific to the sampling-based realisation of iDP.

*d) Membership inference attacks.*: The MIA advantage as we use it follows the formalisation of [22], who introduced the likelihood-ratio attack (LiRA) as a tight empirical estimator. [23] document the privacy onion effect: not every data point is equally vulnerable at the same nominal budget. Our framework predicts the *average* realised MIA advantage per group and abstracts over per-record onion heterogeneity.

*e) Federated learning incentive design.*: A separate literature studies privacy-aware incentive design in federated learning: a server commits to incentives, and heterogeneous clients respond with their own privacy choices, often modelled as a Stackelberg game [24]–[27]. Our framework provides a foundational baseline on which such incentive schemes must rest: it characterises the equilibrium that emerges when no incentive is offered, which is what any incentive design needs to improve upon.

*f) Utility-noise tradeoff in DP-SGD.*: We motivate the functional form of the model-utility value  $V(\sigma)$  via the excess-empirical-risk bound of [7] for convex, Lipschitz DP-ERM, who prove that the excess risk of DP-SGD scales as  $O(d/(n\epsilon)^2)$  at fixed  $\delta$ , equivalently as  $O(\sigma^2/n^2)$  under the moments accountant of [10]. The convex setting is restrictive; for deep learning, [15] provide extensive empirical evidence that test accuracy degrades monotonically with  $\sigma$  on standard benchmarks. We use  $V(\sigma) = 1/(1+\sigma^2)$  as a directional proxy; the qualitative results of our paper do not depend on the precise functional form.

## VII. DISCUSSION AND CONCLUSION

We have shown that sampling-based individual differential privacy is not just *vulnerable* to collusion attacks but *equilibrium-asymmetric*. The pure Nash equilibria of the

Privacy Externality Game, when they exist on the dense action grid, sit *off the diagonal*: one player accepts a higher realised MIA advantage and is compensated by a higher model-utility benefit; the other absorbs the opposite trade. The budget-manipulation attack of [5] is precisely the asymmetric equilibrium played by a coalition that weights the target’s harm in its joint utility, and Theorem 9 shows this is the rational best response of *any* coordinated subset of data holders.

*a) Limitations of this work.*: Three caveats are worth noting. First, our model assumes one-shot play. Real deployments of iDP are dynamic: contributors join and leave; budgets are re-negotiated; trainers re-train. A repeated-game extension is the natural next step, since the asymmetric equilibrium we observe could in principle be averaged out across rounds (or exacerbated, as a player who is repeatedly assigned the high- $A$  role would exit the system). Second, we assume full information about other players’ actions. Real contracts may keep budgets private, and Theorem 6 already distinguishes between fully aware players and players who only see their nominal  $(\epsilon_i, \delta)$  contract. A Bayesian PEG that treats  $\epsilon_i$  as a private type drawn from a known distribution is a natural and important extension. Third, our analysis is restricted to the sampling-based realisation of iDP. The sensitivity-based realisation [2] adjusts per-sample clipping bounds rather than per-sample sampling rates and does not exhibit the same externality (the noise-to-sensitivity ratio is preserved by construction). Our framework specialises to that case as the trivial PEG in which the externality vanishes.

*b) The role of the model-utility function  $V$ .*: Our choice  $V(\sigma) = 1/(1+\sigma^2)$  is motivated by the convex DP-ERM excess-risk bound of [7], and we use it as a directional proxy in our possibly non-convex setting. The qualitative findings—existence of asymmetric NEs, monotone coalition attack, externality lemma—do not depend on the precise shape of  $V$ ; they require only that  $V$  is continuous, monotone-decreasing in  $\sigma$ , and bounded. We expect the same qualitative picture under any utility function with these three properties.

*c) Mitigation.*: The natural mitigation is to make the contract internalise the externality. The  $(\epsilon_i, \delta, \Delta)$ -iDP formulation of [9] adds a  $\Delta$ -divergence bound on the realised vulnerability, restricting the action space of the PEG to profiles that satisfy a per-record bound on excess advantage. We expect this restriction to admit a unique equilibrium and to drive the realised advantage gap between players to within  $\Delta$  by definition. A clean characterisation of the constrained PEG is left to future work.

*d) Open questions.*: The following questions naturally arise from this work but lie outside its scope:

- 1) Does a pure-strategy NE always survive on the continuous action space under the coupled cost, and does a  $\Delta$ -bound on the slack suffice for existence?
- 2) Is the PEG a potential game in the sense of [17]? PoP’s analysis was a potential game because the privacy cost was separable; the externality term breaks this directly. Conditions on the cross-partial of  $A_i$  that recover the potential structure would unify the two settings.






Fig. 4: **C5.** Target’s realised MIA advantage as a function of the other group’s budget  $\varepsilon_O$  and proportion, at two fixed target budgets  $\varepsilon_T \in \{4, 16\}$ . The curves cross at  $\varepsilon_O = \varepsilon_T$ . Below crossover, larger proportions raise  $A_t$ ; above crossover, larger proportions lower it.

- 3) What is the sign of the externality across composition regimes (minority vs majority weak-privacy)? Our experiments suggest the smaller group is favoured, but the analytic dependence on  $|G_g|$  deserves a closed-form characterisation.
- 4) What is the appropriate player granularity (single sample versus silo) in the FL setting, and how does the granularity change the equilibria?
- 5) Worst-case (analytical) vs realised (empirical) cost—which should the contract be written against, and how large is the gap in practice across deep-learning workloads?
- 6) Is the  $(\varepsilon_i, \delta, \Delta)$ -iDP contract by itself sufficient as a collusion-resistant mechanism, or are sensitivity-based iDP or public-data padding needed when the slack is exceeded?

*e) Wider lesson.:* The wider lesson is that an *individual* privacy contract is incomplete: it does not pin down the collective dynamics. Where individuals interact strategically (in a single machine-learning training, in a federated round, in any future composition of the iDP mechanism), the contract must also constrain their joint behaviour or the equilibrium will quietly violate at least one participant’s expectations while preserving every formal guarantee.

*f) Conclusion.:* The promise of iDP—that every data contributor controls her own privacy—is broken by the privacy externality of sampling-based mechanisms. We formalised the strategic situation faced by iDP data contributors as the Privacy Externality Game, proved monotonicity of the externality and characterised the equilibria. The result is that the collusion attack on iDP is not adversarial but *rational*, and the remedy must operate at the level of the contract.

## REFERENCES

- [1] C. Dwork, F. McSherry, K. Nissim, and A. Smith, “Calibrating noise to sensitivity in private data analysis,” in *Theory of Cryptography (TCC)*, 2006, pp. 265–284.
- [2] F. Boenisch, C. Mühl, A. Dziedzic, R. Rinberg, and N. Papernot, “Have it your way: Individualized privacy assignment for DP-SGD,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [3] Z. Jorgensen, T. Yu, and G. Cormode, “Conservative or liberal? personalized differential privacy,” in *IEEE International Conference on Data Engineering (ICDE)*, 2015, pp. 1023–1034.
- [4] V. Feldman and T. Zrnic, “Individual privacy accounting via a Rényi filter,” *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- [5] J. Kaiser, A. Ziller, E. Triantafyllou, D. Rückert, and G. Kaissis, “Your privacy depends on others: Collusion vulnerabilities in individual differential privacy,” in *IEEE Conference on Secure and Trustworthy Machine Learning (SaTML)*, 2026.
- [6] B. Pejó, Q. Tang, and G. Biczók, “Together or alone: The price of privacy in collaborative learning,” in *Proceedings on Privacy Enhancing Technologies (PoPETs)*, vol. 2019, no. 2, 2019, pp. 47–65.
- [7] R. Bassily, A. Smith, and A. Thakurta, “Private empirical risk minimization: Efficient algorithms and tight error bounds,” in *IEEE 55th Annual Symposium on Foundations of Computer Science (FOCS)*, 2014, pp. 464–473.
- [8] B. Balle, G. Barthe, and M. Gaboardi, “Privacy profiles and amplification by subsampling,” *Journal of Privacy and Confidentiality*, vol. 10, no. 1, 2020.
- [9] G. Kaissis, S. Kolek, B. Balle, J. Hayes, and D. Rückert, “Beyond the calibration point: Mechanism comparison in differential privacy,” in *International Conference on Machine Learning (ICML)*, 2024.
- [10] M. Abadi, A. Chu, I. Goodfellow, H. B. McMahan, I. Mironov, K. Talwar, and L. Zhang, “Deep learning with differential privacy,” in *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, 2016, pp. 308–318.
- [11] J. Dong, A. Roth, and W. J. Su, “Gaussian differential privacy,” *Journal of the Royal Statistical Society Series B*, vol. 84, no. 1, pp. 3–37, 2022.
- [12] Google, “Google’s Differential Privacy Accounting Library,” <https://github.com/google/differential-privacy>, 2023.
- [13] J. F. Nash, “Equilibrium points in  $n$ -person games,” *Proceedings of the National Academy of Sciences*, vol. 36, no. 1, pp. 48–49, 1950.
- [14] E. Koutsoupias and C. Papadimitriou, “Worst-case equilibria,” in *Symposium on Theoretical Aspects of Computer Science (STACS)*, 1999, pp. 404–413.

- [15] F. Tramèr and D. Boneh, “Differentially private learning needs better features (or much more data),” in *International Conference on Learning Representations (ICLR)*, 2021.
- [16] J. Hayes, S. Mahloujifar, N. Carlini, and F. Tramer, “Bullseye polytope: A scalable clean-label poisoning attack with improved transferability,” in *IEEE Symposium on Security and Privacy (S&P)*, 2024.
- [17] D. Monderer and L. S. Shapley, “Potential games,” *Games and Economic Behavior*, vol. 14, no. 1, pp. 124–143, 1996.
- [18] I. Mironov, “Rényi differential privacy,” in *IEEE 30th Computer Security Foundations Symposium (CSF)*, 2017, pp. 263–275.
- [19] S. Ioannidis and P. Loiseau, “Linear regression as a non-cooperative game,” in *International Conference on Web and Internet Economics (WINE)*, 2013, pp. 277–290.
- [20] M. Chessa, J. Grossklags, and P. Loiseau, “A game-theoretic study on non-monetary incentives in data analytics projects with privacy implications,” in *IEEE Computer Security Foundations Symposium (CSF)*, 2015, pp. 90–104.
- [21] J. Pawlick and Q. Zhu, “A Stackelberg game perspective on the conflict between machine learning and data obfuscation,” in *IEEE International Workshop on Information Forensics and Security (WIFS)*, 2016.
- [22] N. Carlini, S. Chien, M. Nasr, S. Song, A. Terzis, and F. Tramer, “Membership inference attacks from first principles,” in *IEEE Symposium on Security and Privacy (S&P)*, 2022, pp. 1897–1914.
- [23] N. Carlini, M. Jagielski, C. Zhang, N. Papernot, A. Terzis, and F. Tramer, “The privacy onion effect: Memorization is relative,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [24] Z. Li, Z. Lin, A. Yi, and X. Chen, “A game-theoretic framework for privacy-aware client sampling in federated learning,” *arXiv preprint arXiv:2412.05636*, 2024.
- [25] K. Donahue and J. Kleinberg, “Optimality and stability in federated learning: A game-theoretic approach,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- [26] e. a. Tu, “ID-HFL: Incentive-driven heterogeneous federated learning based on game theory and differential privacy,” *Knowledge-Based Systems*, 2025.
- [27] L. Ma *et al.*, “Differential privacy in federated learning: An evolutionary game analysis,” *Applied Sciences*, vol. 15, no. 6, p. 2914, 2025.

## APPENDIX A

### DETAILED MECHANISM-SOLVER CONSTRUCTION

We elaborate on the bisection used to obtain the mechanism solution  $(\sigma, p_1, \dots, p_n)$  for an action profile  $\varepsilon$ .

*a) Inner step:  $p_g$  from  $(\sigma, \varepsilon_g)$ .* For each fixed  $\sigma$  and each group  $g$ , we want the largest  $p_g \in [0, 1]$  satisfying the  $I$ -fold  $(p_g, \sigma)$ -subsampled Gaussian at  $(\varepsilon_g, \delta)$ -DP under (C1). Since the  $\varepsilon$  value of the subsampled Gaussian is strictly increasing in  $p_g$  (more sampling  $\Rightarrow$  more privacy loss), a unique  $p_g \in (0, 1]$  exists, capped at 1 if the budget cannot be exhausted. We compute the per-group  $\varepsilon$  by Rényi-DP composition [18] and binary-search  $p_g$  to bring this  $\varepsilon$  to  $\varepsilon_g$  to within a relative tolerance of  $10^{-3}$ .

*b) Outer step:  $\sigma$  from  $(p_g, \varepsilon)$ .* Define  $B(\sigma) = \sum_g |G_g| p_g(\sigma, \varepsilon)$  as the realised expected batch size. Since each  $p_g$  is monotone-increasing in  $\sigma$ , so is  $B$ . We bisect  $\sigma$  to bring  $B$  to the target  $E_b$ , again to within a relative tolerance of  $10^{-3}$ .

*c) Participation threshold.* We declare profile infeasible (per Theorem 4) if any  $p_g$  falls below  $10^{-5}$ . Below this threshold the group’s contribution to (C2) is numerically negligible and the iDP contract is trivially satisfied. We verified the qualitative results are insensitive to the choice of this threshold across  $[10^{-6}, 10^{-3}]$ .

*d) MIA advantage.* Given the mechanism solution  $(\sigma, p_g)$ , we compute the MIA advantage of group  $g$  exactly from the privacy-loss-distribution of the  $I$ -fold  $(p_g, \sigma)$ -subsampled

Gaussian using [12], then evaluate the trade-off function and Eq. (2). This is the same formulation used by [5] and yields the tightest analytical MIA advantage for the resulting mechanism.

## APPENDIX B

VALIDITY OF  $V(\sigma) = 1/(1 + \sigma^2)$

*a) Frame.* The convex-DP-ERM excess-risk bound of [7] is, for a 1-Lipschitz  $L$ -strongly-convex  $\ell: \mathcal{X} \times \mathcal{D} \rightarrow \mathbb{R}$ ,

$$\mathbb{E}[\hat{L}(\hat{\theta}) - \hat{L}(\theta^*)] = O\left(\frac{d \log(1/\delta)}{n^2 \varepsilon^2}\right).$$

Translating to noise multiplier via the moments accountant of [10], this reads  $O(\sigma^2 \cdot d/n^2)$  in excess empirical risk. Treating “model utility” as “excess risk” and re-normalising to  $V \in [0, 1]$  produces  $V \propto 1/(1 + \alpha \sigma^2)$  for a problem-dependent  $\alpha$ . We set  $\alpha = 1$  for simplicity; sensitivity analysis (Section C) shows the qualitative results are unchanged for  $\alpha \in [0.1, 10]$ .

*b) Beyond convex.* For non-convex deep-learning losses no analogous closed-form bound is known. The empirical literature [2], [10], [15] reports monotone test-accuracy degradation with  $\sigma$  in standard benchmarks. We use  $V(\sigma) = 1/(1 + \sigma^2)$  as a *directional* proxy; the qualitative results of our paper (existence of asymmetric NEs, monotone coalition attack, externality lemma) do not depend on the precise shape of  $V$ ; they require only that  $V$  is continuous, monotone-decreasing in  $\sigma$ , and bounded.

## APPENDIX C

### SENSITIVITY TO GRID RESOLUTION AND TO $V$

*a) Grid refinement.* For the 2-player setting of Fig. 1 we re-ran the sweep with  $K = 44$  log-spaced points (twice the resolution) and verified that the set of pure NEs is contained, modulo the new intermediate grid points, in the set found at  $K = 22$ . Specifically, the asymmetric NEs at  $K = 22$  persist at  $K = 44$  and the new intermediate grid points do not introduce additional NEs.

*b) Functional form of  $V$ .* The same 2-player sweep was repeated with  $V(\sigma) = e^{-\sigma}$  (an exponential decay alternative); the pure NEs shift quantitatively (mostly toward higher  $\varepsilon$  values, since  $V_{\text{exp}}$  falls off faster than  $V_{1/(1+\sigma^2)}$  as  $\sigma$  grows) but the asymmetric structure is preserved.

*c) Participation threshold.* Lowering the threshold to  $10^{-6}$  extends the feasible region; the NEs found at  $10^{-5}$  remain NEs but additional ones appear in the newly feasible interior. Raising the threshold to  $10^{-3}$  shrinks the feasible region and removes some NEs at the boundary; the surviving NEs do not move.